

SIMATS ENGINEERING

CO3-ASSESSMENT TOOL2-MCQ

COURSE NAME: Artificial Intelligence

COURSE CODE: CSA17 09

COURSE OUTCOME COVERED

CO3: Analyze knowledge representation and reasoning using propositional logic, FOL, inference, resolution, chaining methods, knowledge representation to perform logical reasoning for drawing valid conclusions and decision making. (BL4)

Assessment Tool Weightage: 10%

QUESTIONS: 30

Total Marks:30

Code of Conduct:

I _DHANUSH.G (192512132)_ (Name / Reg No) certify that this submission is my original work and that I have adhered to the guidelines specified for this assessment.
I understand that any violation of academic integrity rules will result in disciplinary action.

Signature of the student: Dhanush.G

Circle the most appropriate option for each question.

Q1. Which of the following best describes a Logical Agent in AI?

- | | |
|--|---|
| a) An agent that reacts randomly to the environment | c) An agent that only follows pre-defined scripts |
| b) An agent that uses a knowledge base and logical to decide actions | d) An agent that learns from rewards inference |

Q2. In Propositional Logic, 'P $\rightarrow$ Q' is logically equivalent to:

- a) $P \wedge \neg Q$ c) $\neg Q \rightarrow \neg P$ only
b) $\neg P \vee Q$ d) Both b and c

Q3. Conjunctive Normal Form (CNF) requires the formula to be expressed as:

- a) A disjunction of conjunctions c) A conjunction of literals only
b) A conjunction of disjunctions (clauses) d) A sequence of implications

Q4. In FOL, the sentence $\forall x: \text{Student}(x) \rightarrow \text{Passes}(x)$ means:

- a) Some students pass c) If something passes, it is a student
b) Every entity that is a student passes d) There exists at least one student who passes

Q5. Unification in FOL is the process of finding substitutions that:

- a) Convert clauses to CNF c) Prove a goal using backward chaining
b) Make two logical expressions syntactically identical d) Derive facts using Modus Ponens

Q6. In the Resolution algorithm, deriving the empty clause $\square$ indicates:

- a) The knowledge base is consistent c) A contradiction — the original goal is proven
b) The negated goal is satisfiable d) More clauses need to be added

Q7. Forward Chaining is best described as a _____ strategy.

- a) Goal-driven / top-down c) Heuristic-based depth-first
b) Data-driven / bottom-up d) Probabilistic inference

Q8. Backward Chaining begins reasoning from:

- a) All known facts toward a conclusion c) A random clause in the knowledge base
b) The goal and works backward to verify supporting facts d) The resolution refutation tree

Q9. Which step is NOT part of converting a sentence to CNF?

- a) Eliminate biconditionals ($\leftrightarrow$) c) Apply Skolemization to remove existential quantifiers
b) Move negations inward using De Morgan's laws d) Distribute AND over OR

Q10. Knowledge Engineering in FOL involves:

- a) Training a neural network on domain-specific data c) Writing heuristic search strategies
b) Defining domain concepts, relations, and axioms formally in FOL d) Encoding reward functions for reinforcement learning

Q11. Modus Ponens states: If $P \rightarrow Q$ is known and P is true, then:

- a) $\neg Q$ is true c) P is uncertain
b) Q must be true d) The KB is inconsistent

Q12. The Most General Unifier (MGU) is the substitution that:

- a) Is the most specific possible match c) Eliminates all variables from a clause
b) Makes two terms identical with minimal variable d) Applies resolution to produce CNF binding

Q13. Which inference method is more appropriate when a specific goal is known?

- a) Forward Chaining c) Backward Chaining
b) Truth Table method d) A* search

Q14. Propositional Logic differs from First Order Logic because it:

- a) Supports probabilistic reasoning c) Is always more expressive
- b) Cannot represent objects, relations, or quantifiers — d) Handles uncertainty better only truth values of atomic propositions

Q15. Resolution Refutation proves a goal G by:

- a) Directly deriving G from the KB using forward chaining c) Checking G in the truth table
- b) Adding $\neg G$ to the KB and deriving a contradiction d) Applying Modus Tollens repeatedly (empty clause $\square$)

Q16. Learning from observation in AI refers to:

- a) An agent receiving numerical rewards from the environment c) Acquiring knowledge by following explicit symbolic rules
- b) Improving performance by examining examples or experience from the environment d) Memorizing a fixed dataset without generalization

Q17. Inductive learning is best described as:

- a) Deriving specific conclusions from general rules c) Memorizing training data exactly
- b) Generalizing a hypothesis from specific training examples d) Applying search algorithms to find optimal policies

Q18. In a Decision Tree, what does each internal node represent?

- a) A class label / output decision c) A test on an attribute / feature
- b) A probability distribution over outcomes d) A training data sample

Q19. Explanation-Based Learning (EBL) differs from inductive learning because it:

- a) Requires large amounts of training data to generalize c) Uses prior domain knowledge to explain and generalize a single example
- b) Learns only from numeric data using regression d) Applies unsupervised clustering to find hidden patterns

Q20. In statistical learning, a Naive Bayes classifier assumes:

- a) All features are conditionally dependent on each other given the class
- b) The data must follow a uniform distribution
- c) All features are conditionally independent given the class label
- d) Labels are determined by a reward function

Q21. Learning with complete data (no hidden variables) in a Bayesian network involves:

- a) Using the EM algorithm to estimate missing values
- b) Training a deep neural network with dropout regularization
- c) Directly computing maximum likelihood estimates from observed data
- d) Applying reinforcement signals to update weights

Q22. The Expectation-Maximization (EM) algorithm is used when:

- a) Training data is fully labeled and no variables are hidden
- b) An agent explores an unknown environment using trial and error
- c) The dataset contains hidden (latent) variables that cannot be directly observed
- d) Backpropagation fails to converge in a neural network

Q23. In a Neural Network, the activation function serves to:

- a) Initialize the weights of each neuron to zero
- b) Normalize the input dataset before training
- c) Introduce non-linearity to enable learning of complex patterns
- d) Store previously learned examples in memory

Q24. Which of the following correctly describes backpropagation in a neural network?

- a) Weights are updated in the forward pass using input gradients
- b) Outputs of each layer are stored and reused in later epochs
- c) The error is propagated backward from output to layers to adjust weights
- d) A separate model is trained for each hidden layer independently

Q25. A Support Vector Machine (SVM) — a type of kernel machine — classifies data by:

- a) Minimizing the number of training examples used
- b) Grouping similar data points into clusters using centroids
- c) Finding the maximum-margin hyperplane that separates classes
- d) Updating weights using Q-values from an environment

Q26. The kernel trick in kernel machines (SVMs) allows:

- a) Reducing the number of support vectors needed
- b) Accelerating gradient descent convergence
- c) Operating in a high-dimensional feature space without explicitly computing coordinates
- d) Replacing hidden layers in a deep neural network

Q27. In Reinforcement Learning, the agent aims to maximize:

- a) Classification accuracy on a fixed test set
- b) The number of features selected from the training data
- c) Cumulative reward received over time through interaction with the environment
- d) The size of the neural network hidden layer

Q28. The Q-Learning algorithm in Reinforcement Learning is classified as:

- a) A supervised learning algorithm requiring labeled datasets
- b) An unsupervised clustering technique using centroids classifier
- c) A model-free, off-policy temporal difference learning method
- d) A decision-tree-based ensemble

Q29. Overfitting in machine learning occurs when:

- a) The model performs poorly on both training and test data
- b) The learning rate is set too low during gradient descent
- c) The model fits training data too closely and generalizes poorly to new data
- d) The training set has more features than training examples

Q30. Which of the following best distinguishes Reinforcement Learning from Supervised Learning?

- a) RL uses labeled input-output pairs; SL uses reward signals
 b) Both RL and SL require explicit reward functions
 c) RL learns via trial-and-error with reward feedback; learns from labeled examples
 d) RL is always faster to converge than SL on any dataset

RUBRICS FOR CALCULATION

Criteria	Weightage (%)	Level 4 (Exemplary)	Level 3 (Proficient)	Level 2 (Developing)	Level 1 (Beginning)
Conceptual Understanding	30%	Demonstrates thorough understanding; answers all questions accurately	Shows good understanding with few errors	Partial understanding; several incorrect responses	Lacks understanding; majority of answers incorrect
Accuracy of Answers	25%	All answers are correct	Most answers are correct with minor mistakes	Some correct answers; frequent errors	Mostly incorrect answers
Application of Knowledge	20%	Correctly applies concepts to problem-based questions	Applies concepts with minor errors	Limited ability to apply concepts	Unable to apply concepts
Time Management	15%	Completes quiz well within time efficiently	Completes on time with minor delays	Requires extra time or rushes through questions	Unable to complete within given time
Question Interpretation	10%	Interprets all questions correctly and responds appropriately	Misinterprets a few questions	Often misinterprets questions	Unable to understand questions



CO3-AT-2 MCQ

Total points **30/30** ?

The respondent's email
(192512132.simats@saveetha.com) was recorded
on submission of this form.

Register Number *

192512132

Name *

Dhanush G

✓ Which of the following best describes a Logical Agent in AI? *1/1

- ☐ a) An agent that reacts randomly to the environment
- ☒ b) An agent that uses a knowledge base and logical inference to decide actions ✓
- ☐ c) An agent that only follows pre-defined scripts
- ☐ d) An agent that learns from rewards

